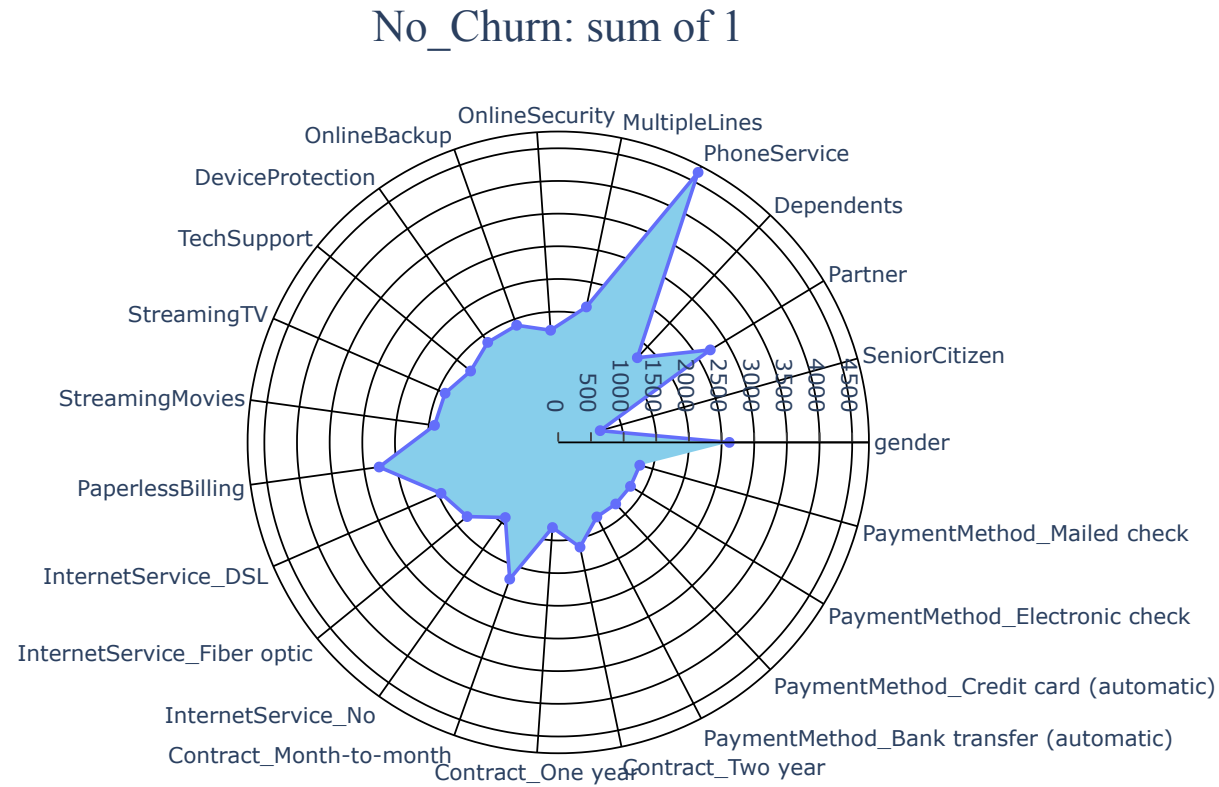


```
In [33]: categories = bifeat_no_churn_radar.keys().tolist()
values = bifeat_no_churn_radar.values.tolist()
radar_chart(categories, values, 'No_Churn: sum of 1')
```



## 2. Model Training Using Grid Search to Find Optimal Parameters and Result Evaluation

```
In [34]: target_col = [LABEL]
df_data = churn_feat_space_and_target
X = df_data.drop(labels = target_col, axis = 1) # pandas dataframe
y = np.where(df_data[LABEL] == LABEL_ONE, 1, 0) # numpy array
```

```

In [35]: def run_ml(X, y, algorithm, algorithm_name, n_folds = None, feature_importance_attr = None):
        """
        Run the machine learning model and show the results.

        Parameters:
        X: input features, dataframe
        y: input labels, numpy array
        algorithm: an object of the model class, e.g., sklearn.linear_model.LogisticRegression
        algorithm_name: name of the model
        n_folds: if not None, use k-fold cross validation, and n_folds is the number of folds.
        feature_importance_attr: if not None, it is the attribute name of the model class to get the feature importance

        Returns:
        algorithm: the trained model
        performance: the model performance dataframe, with columns of 'Model', 'Test Data Size', 'Accuracy', 'Precision (1)',
                     'Recall (1)', 'F1 Score', 'AUC'
        df_feat_importance: the dataframe showing the importance of all features
        """

        # 1. Model training

        if n_folds is not None:
            y_test_pred = y.copy()
            y_test_prob = np.zeros(y.shape)
            # Construct a kfold object
            kf = KFold(n_splits = n_folds, shuffle = True)
            # Iterate through folds
            for train_index, test_index in kf.split(X):
                X_train, X_test = X.iloc[train_index], X.iloc[test_index]
                y_train = y[train_index]
                # Fit the model according to the given training data.
                algorithm.fit(X_train, y_train)
                # Predict class labels for samples in X.
                # Note that the threshold is 0.5
                y_test_pred[test_index] = algorithm.predict(X_test)
                # Return estimates for all classes.
                # predict() will give 0 or 1 as output; predict_proba() will give the probability of 0 (in column 0) and 1 (in column 1).
                y_test_prob[test_index] = algorithm.predict_proba(X_test)[: , 1]
            y_test = y # all data is used as test set
            X_test = X
        else:
            X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.1, random_state = 0)
            # Fit the model according to the given training data.
            algorithm.fit(X_train, y_train)
            # Predict class labels for samples in X.
            # Note that the threshold is 0.5

```

```

y_test_pred = algorithm.predict(X_test)
# Return estimates for all classes.
# predict() will give 0 or 1 as output; predict_proba() will give the probability of 0 (in column 0) and 1 (in column 1).
y_test_prob = algorithm.predict_proba(X_test)[:, 1]

# 2. Show the results

if n_folds is not None:
    print('\nSummary of ' + str(algorithm_name) + ' Model with ' + str(n_folds) + '-Fold Cross Validation:')
else:
    print('\nSummary of ' + str(algorithm_name) + ' Model:')

# 2.1 Print the model details

print('\nAlgorithm:\n')
print(algorithm)

# 2.2 Print the test accuracy

print('\nTest Accuracy:', accuracy_score(y_true = y_test, y_pred = y_test_pred))

# 2.3 Print the classification report

print('\nClassification Report:\n', classification_report(y_true = y_test, y_pred = y_test_pred))

# 2.4 Print and plot the receiver operating characteristic (ROC)

auc = roc_auc_score(y_true = y_test, y_score = y_test_prob)
fpr, tpr, thresholds = roc_curve(y_true = y_test, y_score = y_test_prob)
trace_roc = go.Scatter(x = fpr, y = tpr)
trace_roc_diag = go.Scatter(x = [0, 1], y = [0, 1], line = dict(dash = "dash"))

# 2.5 Calculate and plot the confusion matrix

conf_matrix = confusion_matrix(y_true = y_test, y_pred = y_test_pred)
trace_conf_matrix = go.Heatmap(z = conf_matrix / len(y_test), x = ['Predicted: ' + str(algorithm.classes_[0]),
                                                                'Predicted: ' + str(algorithm.classes_[1])],
                               y = ['True: ' + str(algorithm.classes_[0]), 'True: ' + str(algorithm.classes_[1])])

# 2.6 Calculate and plot accuracy, recall and precision according to different thresholds

if thresholds[0] > 1: # the library code adds a large at index 0 for some specific reason; we don't need it.
    thresholds = thresholds[1:]
accuracy = np.zeros(thresholds.shape)
precision = np.zeros(thresholds.shape)
recall = np.zeros(thresholds.shape)

```

```

for i in range(len(thresholds)):
    y_test_pred_i = (algorithm.predict_proba(X_test)[: , 1] >= thresholds[i])
    cm_i = confusion_matrix(y_true = y_test, y_pred = y_test_pred_i)
    tp = cm_i[1][1]
    tn = cm_i[0][0]
    fp = cm_i[0][1]
    fn = cm_i[1][0]
    accuracy[i] = (tp + tn) / (tp + fp + fn + tn + 0.0)
    recall[i] = tp / (tp + fn + 0.0)
    precision[i] = tp / (tp + fp + 0.0)
trace_acc = go.Scatter(x = thresholds, y = accuracy, name = 'Accuracy')
trace_recall = go.Scatter(x = thresholds, y = recall, name = 'Recall (Positive)')
trace_precision = go.Scatter(x = thresholds, y = precision, name = 'Precision (Positive)')

# 2.7 Calculate and plot feature importance

if feature_importance_attr is not None:
    feature_importance = np.squeeze(getattr(algorithm, feature_importance_attr))
    dict_feat_importance = {'cols' : X.columns, 'importance' : feature_importance}
    df_feat_importance = pd.DataFrame(data = dict_feat_importance)
    df_feat_importance = df_feat_importance.iloc[df_feat_importance.importance.abs().argsort()[::-1]] # [::-1] to get decreasing
order
    trace_feat_importance = go.Bar(x = df_feat_importance['cols'], y = df_feat_importance['importance']) # bar plot

# 3. Show the figures

fig1 = py.subplots.make_subplots(rows = 1, cols = 2, subplot_titles = ('ROC', 'Confusion Matrix'),
                                print_grid = False) # Remove the "This is the format of your plot grid..."
fig1.add_trace(trace_roc, row = 1, col = 1)
fig1.add_trace(trace_roc_diag, row = 1, col = 1)
fig1.add_trace(trace_conf_matrix, row = 1, col = 2)
if n_folds is not None:
    title = str(algorithm_name) + ' (' + str(n_folds) + '-Fold Cross Validation)'
else:
    title = str(algorithm_name)
fig1['layout'].update(title = dict(text = '<b>Performance of ' + title + ' Model</b>', x = 0.5,
                                font = dict(family = 'Times New Roman', size = 25)),
                    showlegend = False,
                    annotations = [dict(x = 0.3, y = 0.1, font = dict(size = 25, color = 'white'),
                                text = 'AUC: ' + str(round(auc, 4)))],
                    plot_bgcolor = 'black')
fig1["layout"]["xaxis1"].update(dict(title = "False Positive Rate", gridcolor = 'grey'))
fig1["layout"]["yaxis1"].update(dict(title = "True Positive Rate", gridcolor = 'grey'))
pyo.iplot(fig1)

fig2 = go.Figure()

```

```

fig2.add_trace(trace_acc)
fig2.add_trace(trace_recall)
fig2.add_trace(trace_precision)
fig2['layout'].update(title = dict(text = '<b>Threshold Plot</b>', x = 0.5), titlefont = dict(size = 20), showlegend = True,
                        plot_bgcolor = "rgb(200, 200, 200)")
fig2["layout"]["xaxis"].update(dict(title = "Threshold"))
fig2["layout"]["yaxis"].update(dict(title = "Score"))
pyo.iplot(fig2)

if feature_importance_attr is not None:
    fig3 = go.Figure()
    fig3.add_trace(trace_feat_importance)
    fig3['layout'].update(title = dict(text = '<b>Feature Importance</b>', x = 0.5), titlefont = dict(size = 20),
                          plot_bgcolor = "rgb(223, 237, 245)")
    fig3["layout"]["xaxis"].update(dict(tickangle = 90, tickfont = dict(size = 10)))
    pyo.iplot(fig3)

# construct the performance dataframe
test_size = X_test.shape[0]
tp = conf_matrix[1][1]
tn = conf_matrix[0][0]
fp = conf_matrix[0][1]
fn = conf_matrix[1][0]
acc = (tp + tn) / (tp + fp + fn + 0.0) # accuracy
rec = tp / (tp + fn + 0.0) # recall
prec = tp / (tp + fp + 0.0) # precision
f1 = 2 * (prec * rec) / (prec + rec)
performance = {'Model' : [algorithm_name + '<br>(N_Folds: ' + str(n_folds) + ')'], 'Test Data Size' : [test_size],
               'Accuracy' : [acc], 'Precision (1)' : [prec], 'Recall (1)' : [rec], 'F1 Score' : [f1], 'AUC' : [auc]}
performance = pd.DataFrame(data = performance)

# construct the feature importance dataframe
feat_importance = None
if feature_importance_attr is not None:
    feat_importance = {'Features' : X.columns, algorithm_name + '<br>(N_Folds: ' + str(n_folds) + ')': feature_importance}
    feat_importance = pd.DataFrame(data = feat_importance)
    feat_importance = feat_importance.set_index(keys = 'Features')

return algorithm, performance, feat_importance

```

```
In [36]: def grid_search(X, y, algorithm, parameters):
        """
        Use grid search to find the optimal paramters for the model
        Parameters:
        X: input features, dataframe
        y: input labels, numpy array
        algorithm: an object of the model class, e.g., sklearn.linear_model.LogisticRegression
        parameters: dict or list of dictionaries. Dictionary with parameters names (string) as keys and lists of parameter settings to
                    try as values, or a list of such dictionaries, in which case the grids spanned by each dictionary in the list are
                    explored. This enables searching over any sequence of parameter settings.

        Returns:
        dictionary of best paramters.
        """

        gs = GridSearchCV(estimator = algorithm, param_grid = parameters, cv = 5, verbose = 1, refit = False)
        gs.fit(X, y)
        print('\nBest Score:', gs.best_score_)
        print('\nBest Parameter Set:', gs.best_params_)
        return gs.best_params_
```

## 2.1 Logistic Regression

```
In [37]: logis = LogisticRegression(max_iter = 100, multi_class = 'ovr', n_jobs = 1, solver = 'liblinear', verbose = 0, warm_start = False)
```

```
In [38]: parameters = {
        'penalty' : ['l1', 'l2'],
        'C' : 0.001 * 10 ** (np.arange(0, 1.02, 0.02) * 4) # 0.001 to 10
    }
    logis_params = grid_search(X, y, logis, parameters)
```

Fitting 5 folds for each of 102 candidates, totalling 510 fits

[Parallel(n\_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

Best Score: 0.8041808873720137

Best Parameter Set: {'C': 4.786300923226385, 'penalty': 'l2'}

[Parallel(n\_jobs=1)]: Done 510 out of 510 | elapsed: 39.4s finished

```
In [39]: logis = LogisticRegression(max_iter = 100, multi_class = 'ovr', n_jobs = 1, C = logis_params['C'],  
                                     penalty = logis_params['penalty'], solver = 'liblinear', verbose = 0, warm_start = False)  
_, performance_lr, feat_importance_lr = run_ml(X, y, logis, 'Logistic Regression', n_folds = None, feature_importance_attr = 'coef_')
```

## Summary of Logistic Regression Model:

### Algorithm:

```
LogisticRegression(C=4.786300923226385, class_weight=None, dual=False,  
                    fit_intercept=True, intercept_scaling=1, l1_ratio=None,  
                    max_iter=100, multi_class='ovr', n_jobs=1, penalty='l2',  
                    random_state=None, solver='liblinear', tol=0.0001, verbose=0,  
                    warm_start=False)
```


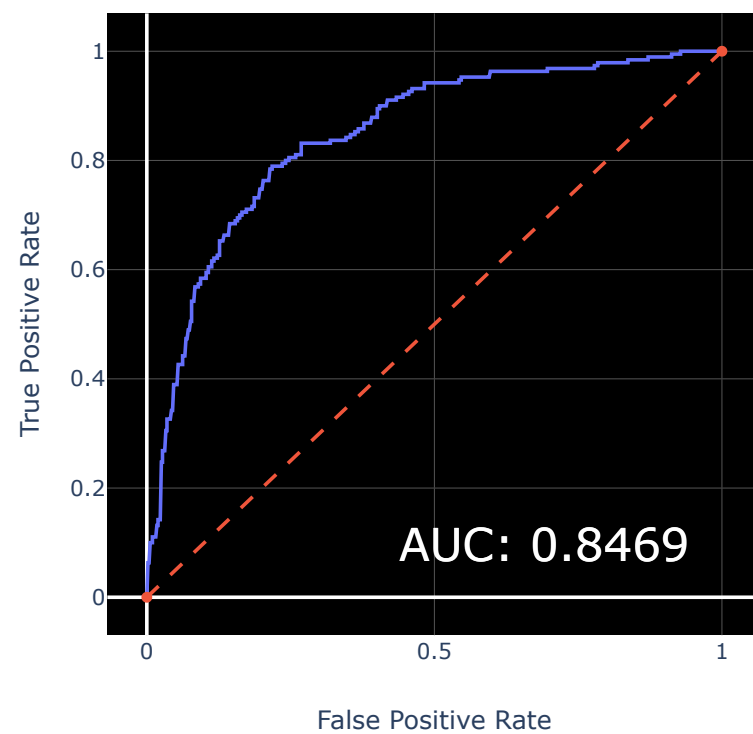
Test Accuracy: 0.8224431818181818

### Classification Report:

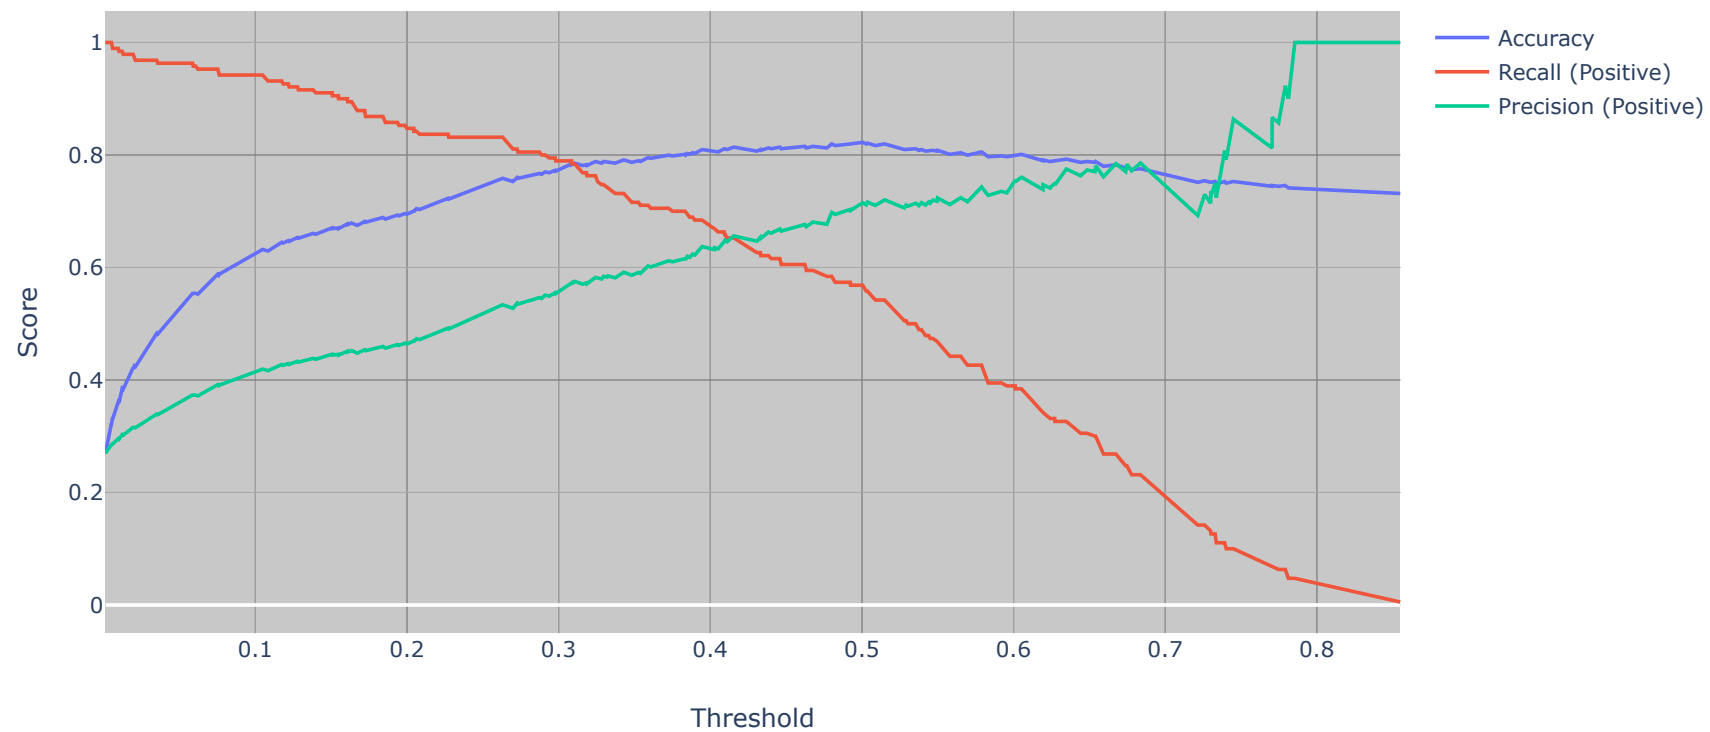
	precision	recall	f1-score	support
0	0.85	0.92	0.88	514
1	0.72	0.57	0.63	190
accuracy			0.82	704
macro avg	0.78	0.74	0.76	704
weighted avg	0.81	0.82	0.82	704




## Performance of Logistic Regression Model



# Threshold Plot



Feature Importance



2.2 Logistic Regression with K-Fold Cross Validation

[illegible]

## Summary of Logistic Regression Model with 5-Fold Cross Validation:

### Algorithm:



```
LogisticRegression(C=4.786300923226385, class_weight=None, dual=False,  
                    fit_intercept=True, intercept_scaling=1, l1_ratio=None,  
                    max_iter=100, multi_class='ovr', n_jobs=1, penalty='l2',  
                    random_state=None, solver='liblinear', tol=0.0001, verbose=0,  
                    warm_start=False)
```

Test Accuracy: 0.8075938566552902


### Classification Report:

	precision	recall	f1-score	support
0	0.85	0.90	0.87	5163
1	0.67	0.56	0.61	1869
accuracy			0.81	7032
macro avg	0.76	0.73	0.74	7032
weighted avg	0.80	0.81	0.80	7032


## Performance of Logistic Regression (5-Fold Cross Validation) Model



# Threshold Plot



## Feature Importance



### 2.3 Logistic Regression with Recursive Feature Elimination (RFE)

The goal of RFE is to select features by recursively considering smaller and smaller sets of features.

[illegible]



```
In [42]: rfe_logis = RFE(logis_rfe, n_features_to_select = n_features_to_select)
```

```
rfe_logis.fit(X, y)
```

```
print("Logistic Regression RFE Result:")
```

```
for k, v in sorted(zip(rfe_logis.ranking_, X.columns)):  
    print(k, ': ', v)
```

```
# get the selected columns
```

```
cols_selected = np.array(X.columns)[rfe_logis.support_].tolist()
```

```
Logistic Regression RFE Result:
```

```
1 : Contract_Month-to-month
```

```
1 : Contract_Two year
```

```
1 : InternetService_Fiber optic
```

```
1 : InternetService_No
```

```
1 : OnlineSecurity
```

```
1 : PaperlessBilling
```

```
1 : PhoneService
```

```
1 : TechSupport
```

```
1 : TotalCharges
```

```
1 : tenure
```

```
2 : PaymentMethod_Mailed check
```

```
3 : PaymentMethod_Credit card (automatic)
```

```
4 : PaymentMethod_Bank transfer (automatic)
```

```
5 : InternetService_DSL
```

```
6 : MultipleLines
```

```
7 : Contract_One year
```

```
8 : StreamingMovies
```

```
9 : SeniorCitizen
```

```
10 : StreamingTV
```

```
11 : OnlineBackup
```

```
12 : Dependents
```

```
13 : MonthlyCharges
```

```
14 : DeviceProtection
```

```
15 : PaymentMethod_Electronic check
```

```
16 : gender
```

```
17 : Partner
```

```
In [43]: X_rfe = X[cols_selected]

_, performance_lr_rfe_kfolds, feat_importance_lr_rfe_kfolds = run_ml(X_rfe, y, logis_rfe, 'Logistic Regression RFE',
                                                                    n_folds = 5, feature_importance_attr = 'coef_')
```

# Summary of Logistic Regression RFE Model with 5-Fold Cross Validation:

## Algorithm:



```
LogisticRegression(C=1, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='ovr', n_jobs=1, penalty='l2', random_state=None,
                    solver='liblinear', tol=0.0001, verbose=0, warm_start=False)
```

Test Accuracy: 0.7987770193401593


## Classification Report:

	precision	recall	f1-score	support
0	0.84	0.89	0.87	5163
1	0.64	0.55	0.59	1869
accuracy			0.80	7032
macro avg	0.74	0.72	0.73	7032
weighted avg	0.79	0.80	0.79	7032


## Performance of Logistic Regression RFE (5-Fold Cross Validation) Model



# Threshold Plot



## Feature Importance



## 2.4 K Nearest Neighbors with K-Fold Cross Validation

```
In [44]: knn = KNeighborsClassifier()
```

```
In [45]: parameters = {  
        'n_neighbors' : np.arange(4, 20, 4)  
        }  
knn_params = grid_search(X, y, knn, parameters)
```

Fitting 5 folds for each of 4 candidates, totalling 20 fits

[Parallel(n\_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

Best Score: 0.7922354948805461

Best Parameter Set: {'n\_neighbors': 16}

[Parallel(n\_jobs=1)]: Done 20 out of 20 | elapsed: 5.3s finished

[illegible]



# Summary of K-Nearest-Neighbors Model with 5-Fold Cross Validation:

## Algorithm:



```
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',  
                    metric_params=None, n_jobs=None, n_neighbors=16, p=2,  
                    weights='uniform')
```

Test Accuracy: 0.7925199089874858

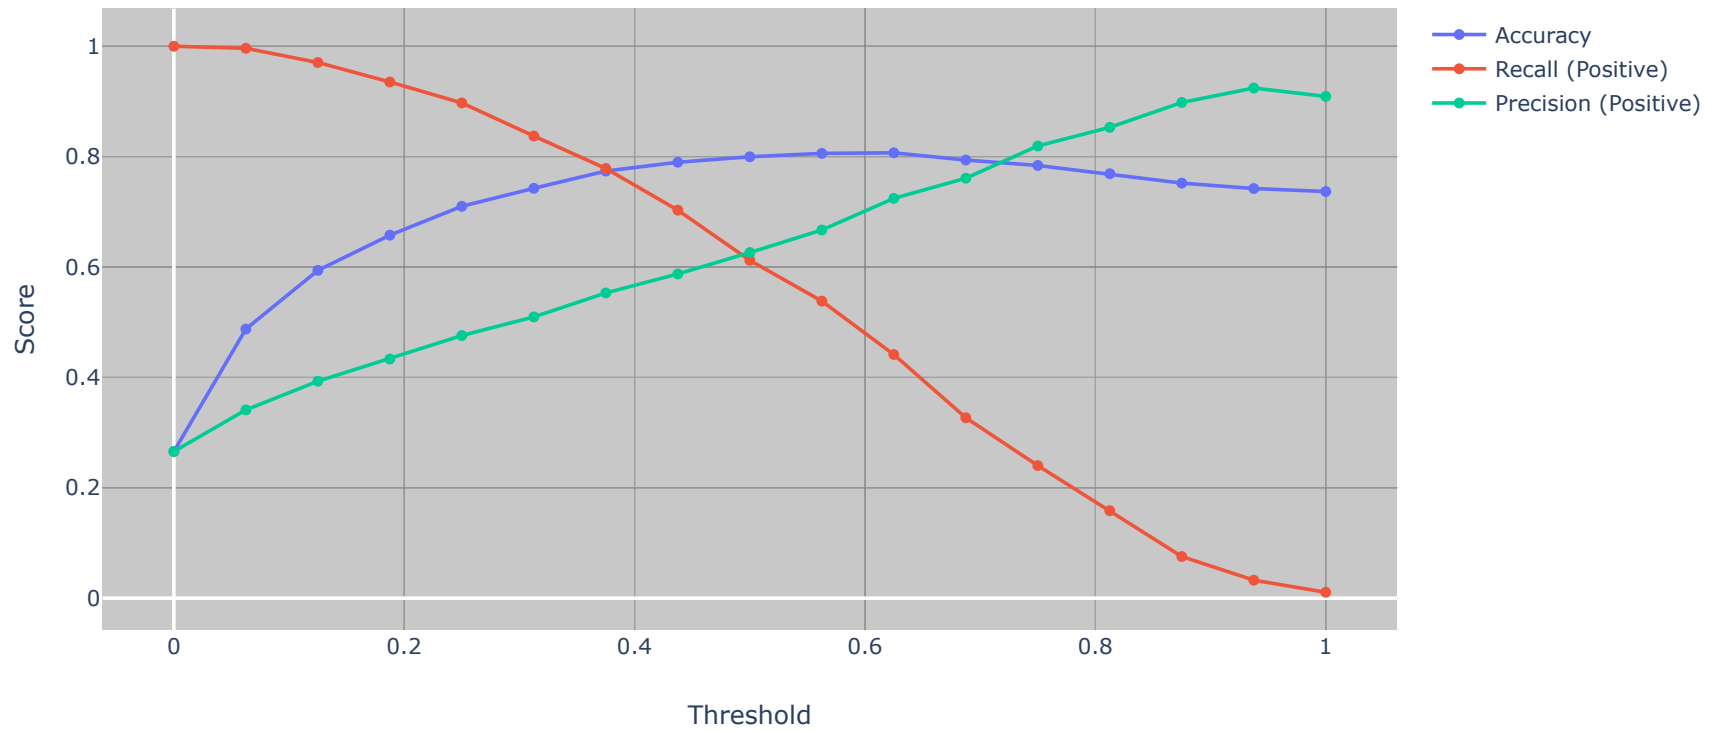
## Classification Report:

	precision	recall	f1-score	support
0	0.83	0.90	0.86	5163
1	0.64	0.51	0.57	1869
accuracy			0.79	7032
macro avg	0.74	0.70	0.71	7032
weighted avg	0.78	0.79	0.78	7032

## Performance of K-Nearest-Neighbors (5-Fold Cross Validation) Model



## Threshold Plot



## 2.5 Random Forest with K-Fold Cross Validation

```
In [47]: rand_forest = RandomForestClassifier(n_estimators = 100, criterion = 'entropy', random_state = 0)
```

```
In [48]: rand_forest, performance_rf_kfolds, feat_importance_rf_kfolds = run_ml(X, y, rand_forest, 'Random Forest', n_folds = 5,  
                                         feature_importance_attr = 'feature_importances_')
```

## Summary of Random Forest Model with 5-Fold Cross Validation:

### Algorithm:



```
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='entropy',  
                        max_depth=None, max_features='auto', max_leaf_nodes=None,  
                        min_impurity_decrease=0.0, min_impurity_split=None,  
                        min_samples_leaf=1, min_samples_split=2,  
                        min_weight_fraction_leaf=0.0, n_estimators=100,  
                        n_jobs=None, oob_score=False, random_state=0, verbose=0,  
                        warm_start=False)
```

Test Accuracy: 0.7864050056882821

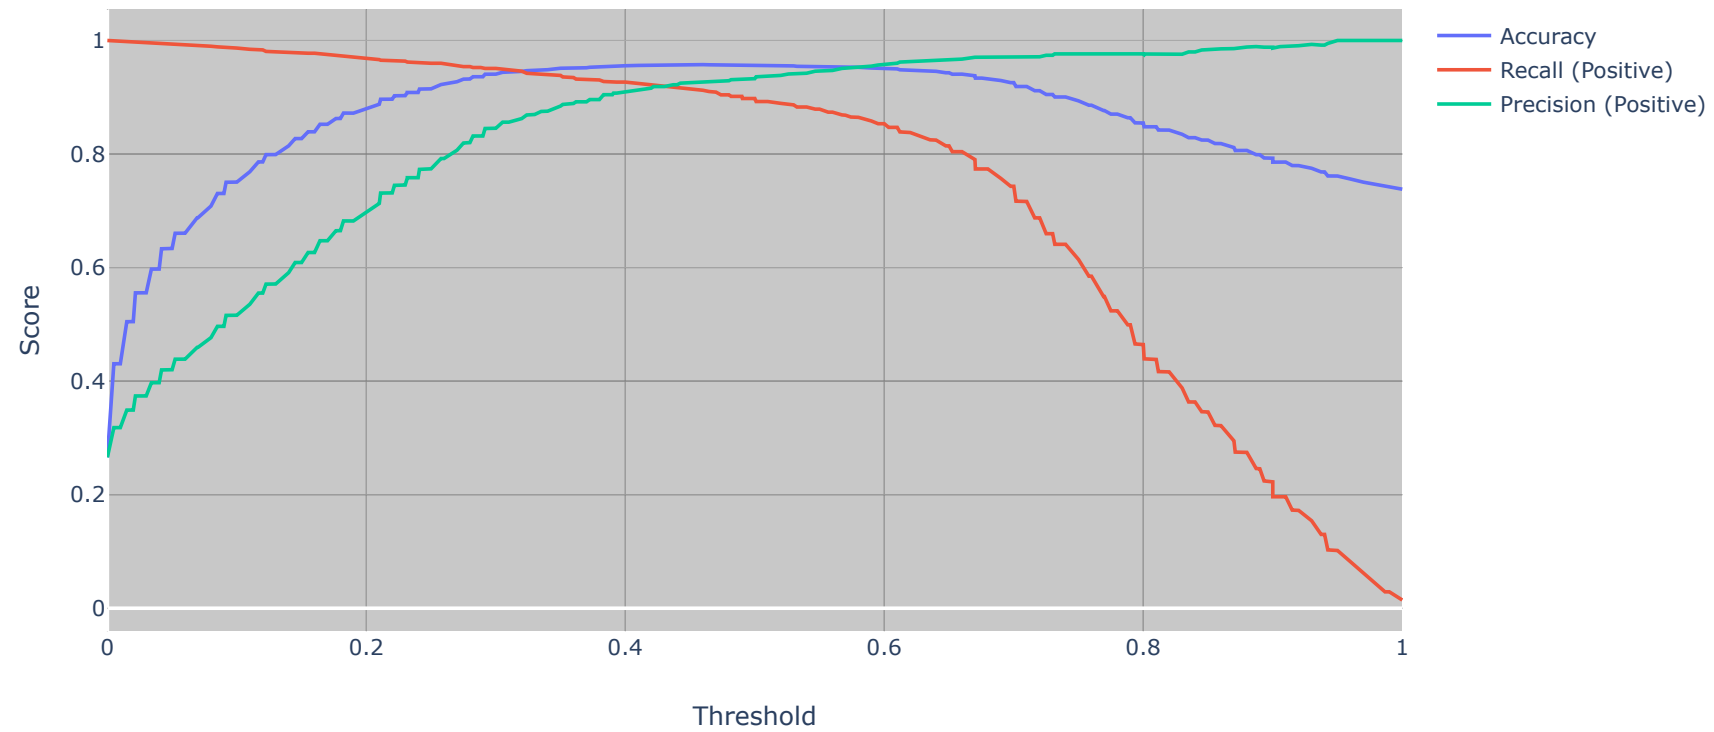
### Classification Report:

	precision	recall	f1-score	support
0	0.83	0.90	0.86	5163
1	0.63	0.49	0.55	1869
accuracy			0.79	7032
macro avg	0.73	0.69	0.70	7032
weighted avg	0.77	0.79	0.78	7032


## Performance of Random Forest (5-Fold Cross Validation) Model



# Threshold Plot



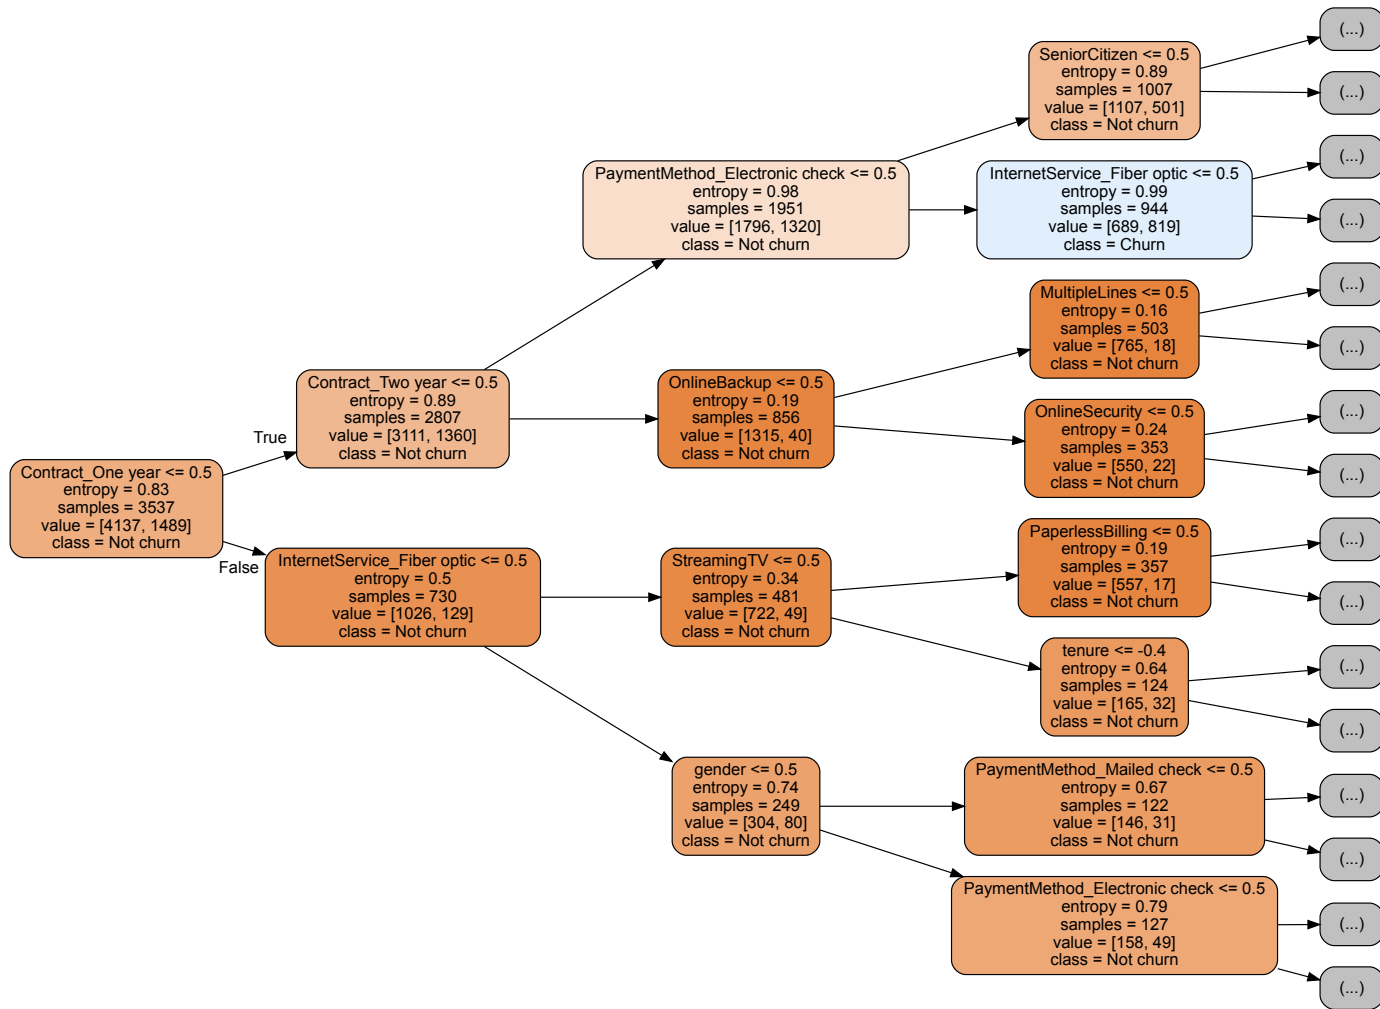
Feature Importance

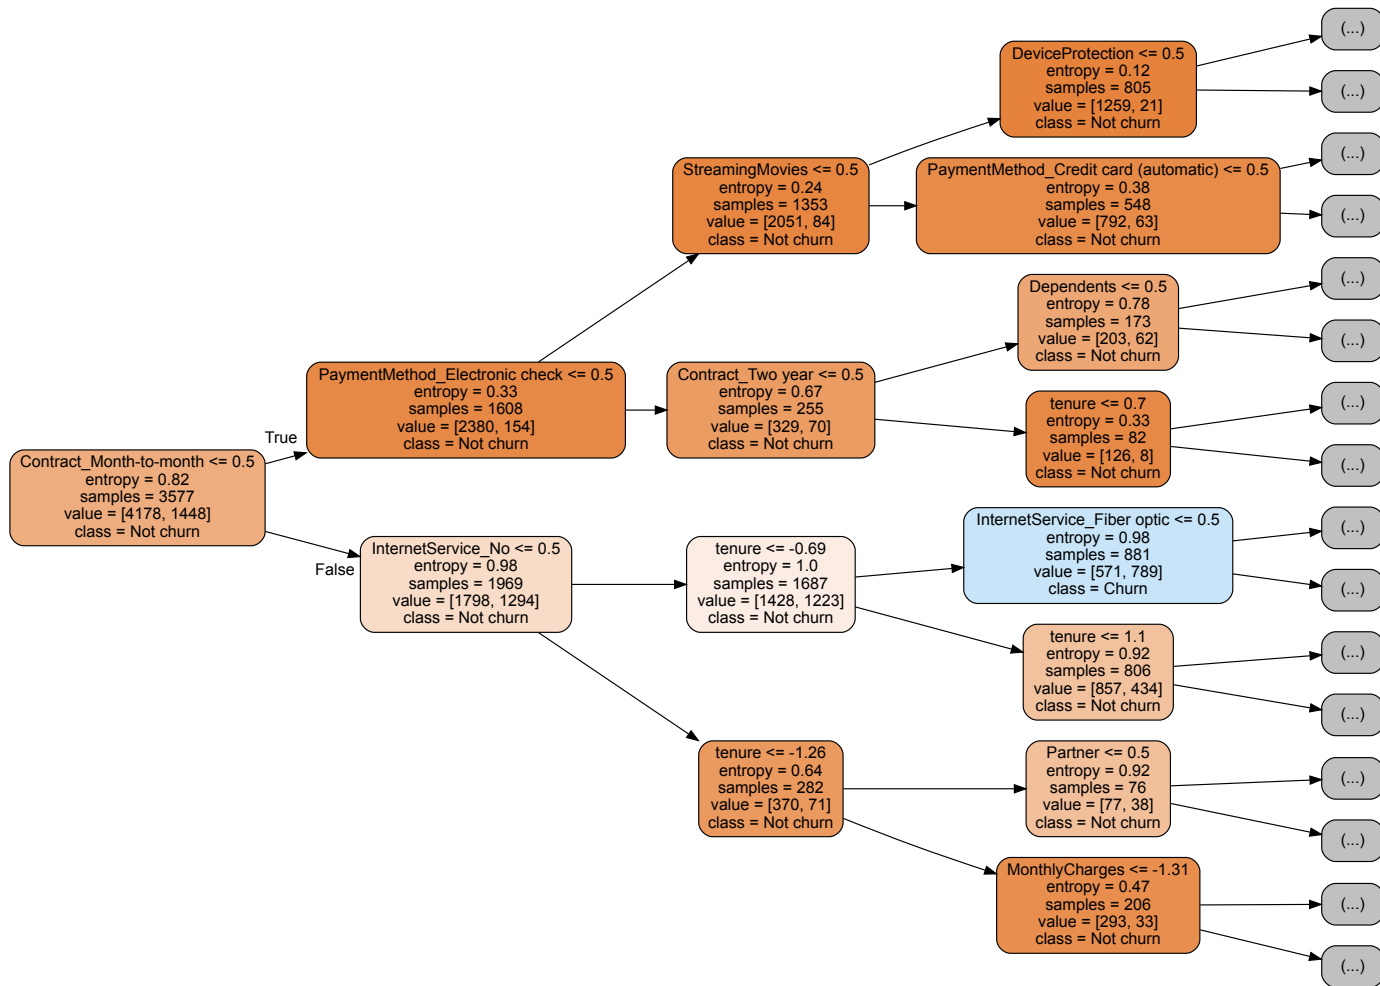


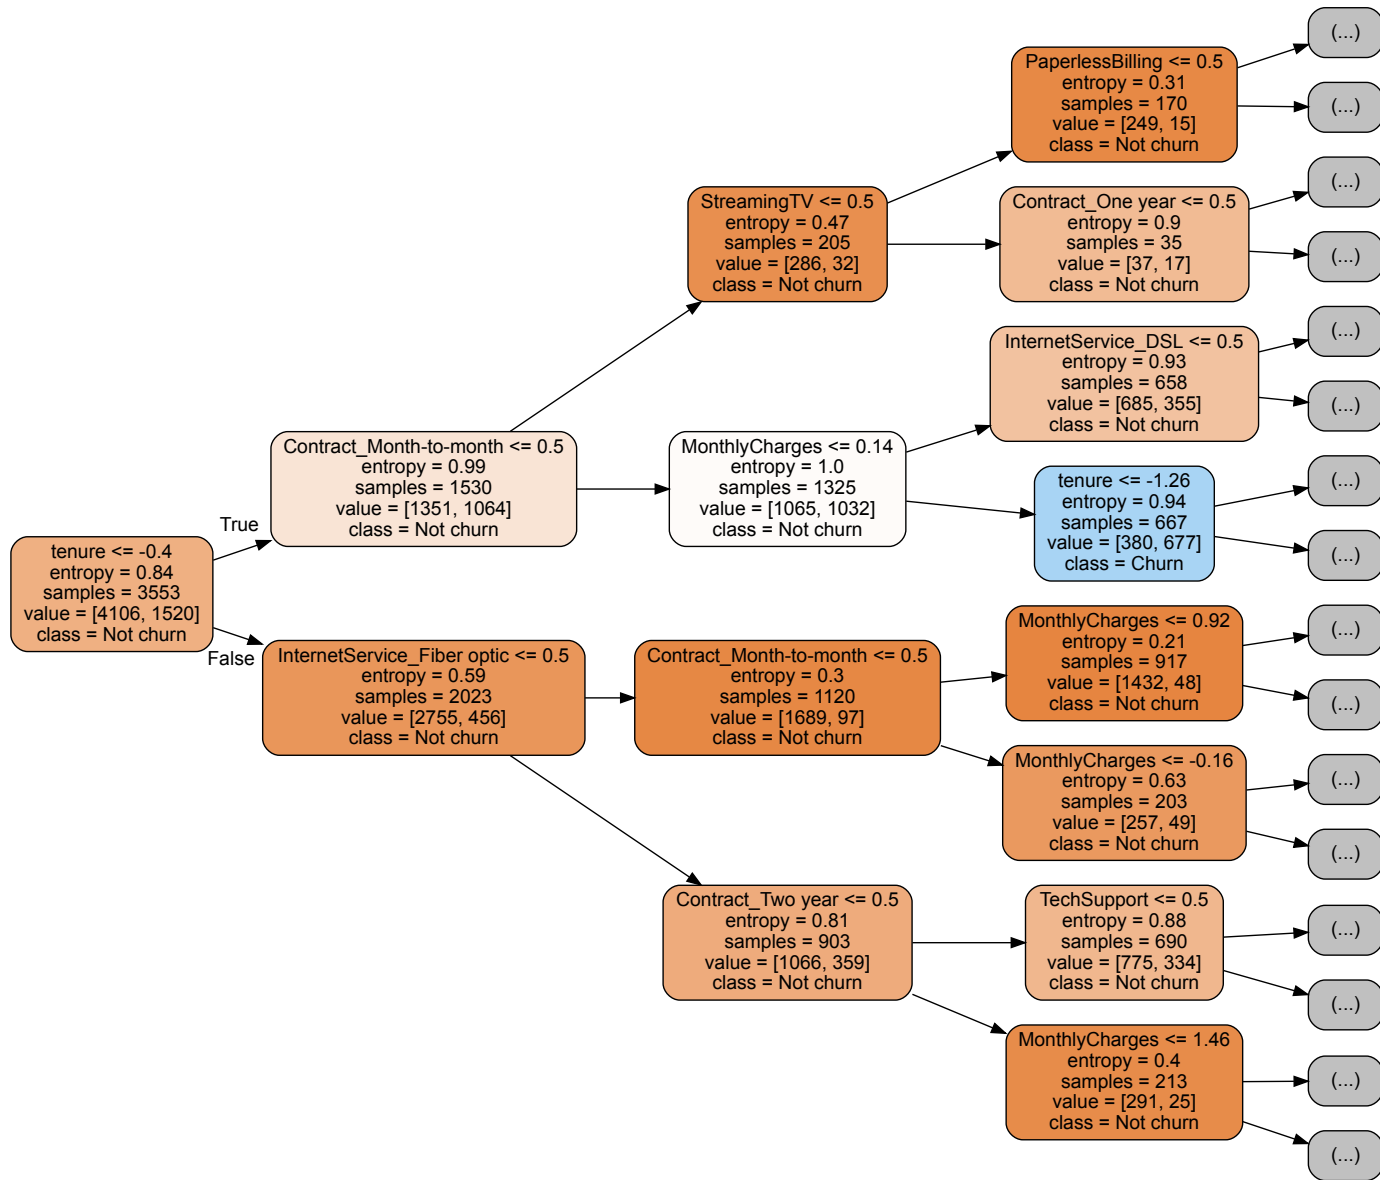


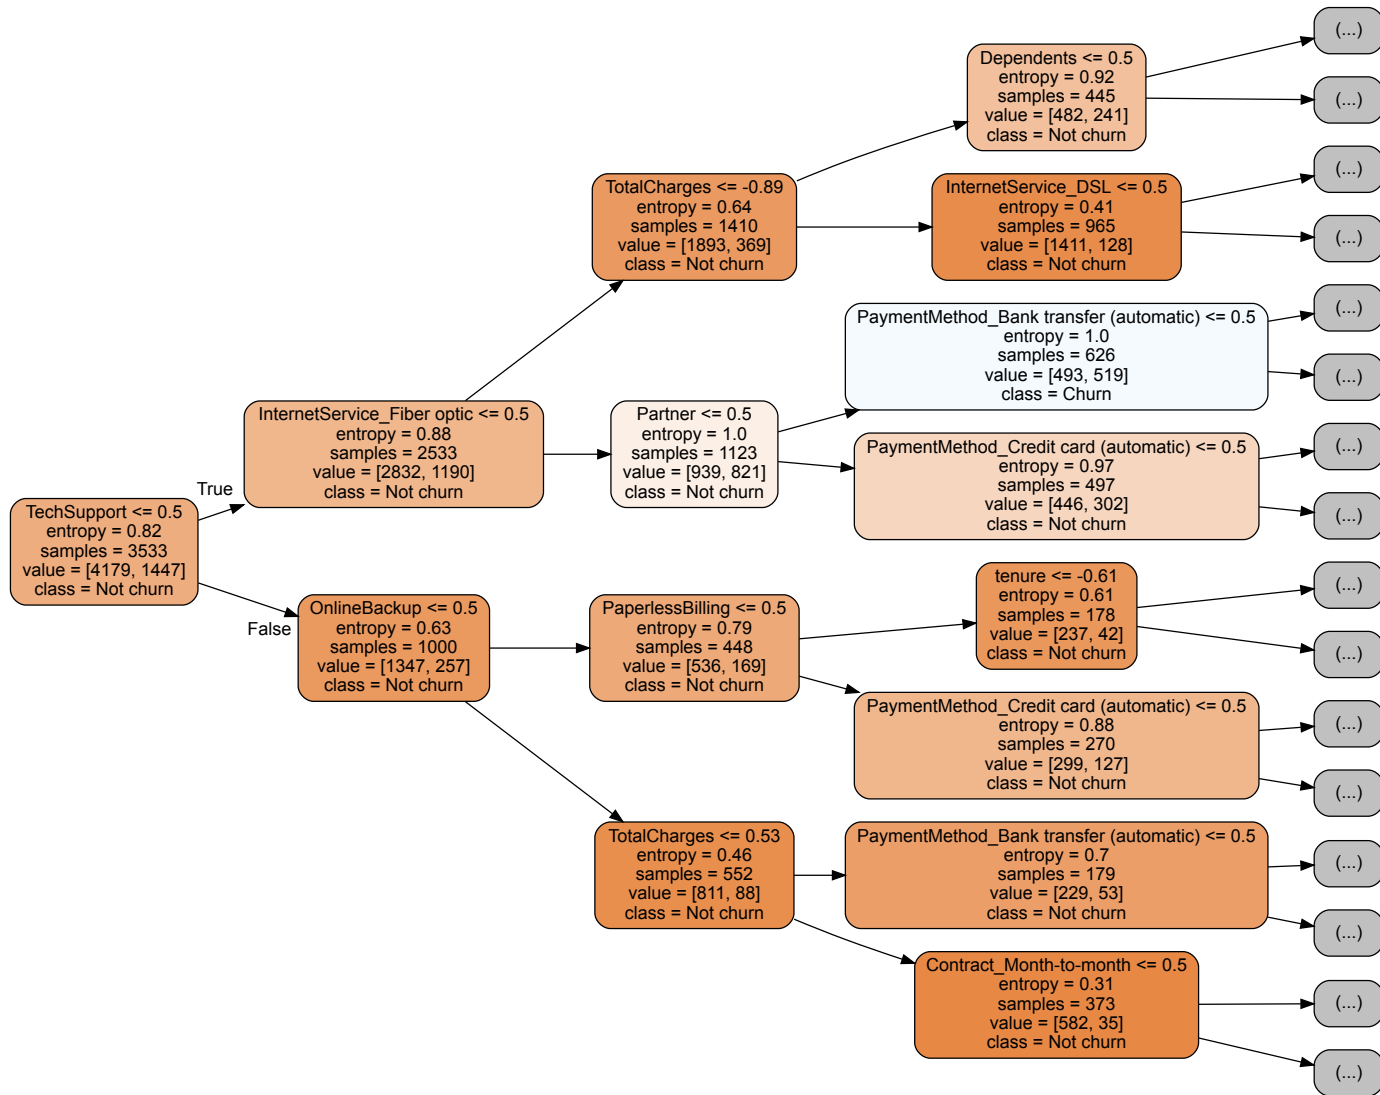
In [49]: *# Plot some of the decision trees*

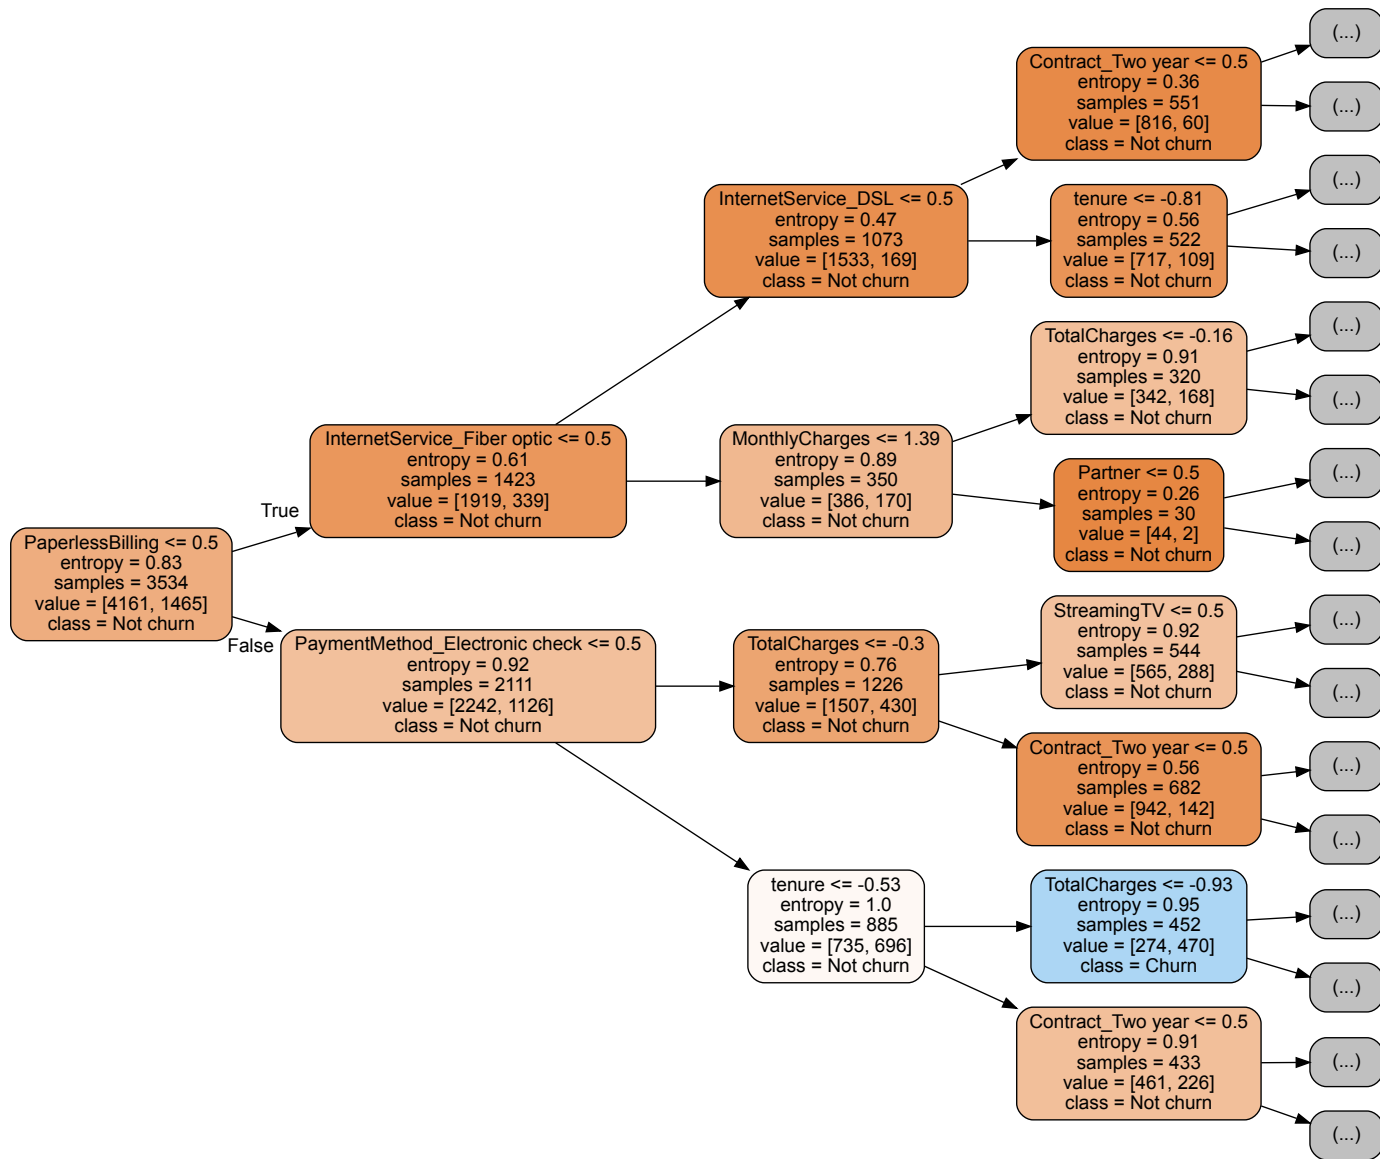
```
for i in range(5):
    dot_data = export_graphviz(rand_forest.estimators_[i], out_file = None, max_depth = 3, feature_names = X.columns.tolist(),
                               rounded = True, proportion = False, precision = 2, class_names = ["Not churn", "Churn"],
                               filled = True, rotate = True)
    pdot = pydot.graph_from_dot_data(dot_data)
    # Access element [0] because graph_from_dot_data actually returns a list of DOT elements.
    pdot[0].set_graph_defaults(size = "\"8,8\"")
    graph = Source(pdot[0].to_string())
    display(graph)
```











**value:** the list represents the count of samples in each class that have reached that node.  
**sample:** if using bootstrap, this number is not equal to the sum of the numbers in 'value'.

## 3. Summary

### 3.1 Model Performance

```
In [50]: # contruct the dataframe table including the performance of all models

performances_df = pd.DataFrame()
performance_df = [performance_lr, performance_lr_kfolds, performance_lr_rfe_kfolds, performance_knn_kfolds, performance_rf_kfolds]
for i in performance_df:
    performances_df = performances_df.append(i, sort = False)
```

In [51]: *# bar chart comparing the performance of different models*


```
metrics = ['Accuracy', 'Precision (1)', 'Recall (1)', 'F1 Score', 'AUC']
traces_performance = []
for metric in metrics:
    trace = go.Bar(y = performances_df['Model'], x = performances_df[metric], orientation = 'h', name = metric)
    traces_performance.append(trace)
layout_performance = go.Layout(title = dict(text = '<b>Model Performance Summary</b>', font = dict(family = 'Times New Roman', size = 25),
                                          x = 0.5),
                              margin = dict(l = 150)) # 'l' in 'margin' sets the left margin in px
fig_performance = go.Figure(data = traces_performance, layout = layout_performance)
pyo.iplot(fig_performance)

# summary table

table_performance = ff.create_table(np.round(performances_df, 4))
pyo.iplot(table_performance)
```



## Model Performance Summary



Model	Test Data Size	Accuracy	Precision (1)	Recall (1)	F1 Score	AUC
Logistic Regression (N_Folds: None)	704	0.8224	0.7152	0.5684	0.6334	0.8469
Logistic Regression (N_Folds: 5)	7032	0.8076	0.6652	0.5559	0.6057	0.8452
Logistic Regression RFE (N_Folds: 5)	7032	0.7988	0.6431	0.5457	0.5904	0.8425
K-Nearest-Neighbors (N_Folds: 5)	7032	0.7925	0.6372	0.5094	0.5662	0.825
Random Forest (N_Folds: 5)	7032	0.7864	0.6268	0.4853	0.547	0.8248


## 3.2 Feature Importance

```
In [52]: feats_importance_df = pd.DataFrame()
feat_importance_df = [feat_importance_lr, feat_importance_lr_kfolds, feat_importance_lr_rfe_kfolds, feat_importance_knn_kfolds,
                      feat_importance_rf_kfolds]
for i in feat_importance_df:
    if i is not None:
        i = i.abs()
        max = i.max() # for normalization, set the max to be 1
        feats_importance_df = pd.concat([feats_importance_df, i / max], axis = 1, sort = False)
```

```
In [53]: traces_feat_importance = []
for i in feats_importance_df.columns:
    trace = go.Bar(y = feats_importance_df.index, x = feats_importance_df[i], orientation = 'h', name = i)
    traces_feat_importance.append(trace)
layout_feat_importance = go.Layout(title = dict(text = '<b>Feature Importance</b>', font = dict(family = 'Times New Roman', size = 25
),
                                   x = 0.5),
                                   yaxis = dict(dtick = 1),
                                   margin = dict(l = 300),
                                   height = 1000)
fig_feat_importance = go.Figure(data = traces_feat_importance, layout = layout_feat_importance)
pyo.iplot(fig_feat_importance)
```

# Feature Importance





```
In [54]: # summary table

print('Feature Importance Summary:')
table_feat_importance_mean = ff.create_table(pd.DataFrame(feats_importance_df.mean(axis = 1).sort_values(axis = 0)[::-1],
                                                         columns = ['Mean Importance']).reset_index())
pyo.ipplot(table_feat_importance_mean)
```

Feature Importance Summary:

Features	Mean Importance
tenure	0.9630173596922257
TotalCharges	0.6457857584345441
InternetService_No	0.6267839269666561
Contract_Two year	0.4957349744170466
MonthlyCharges	0.4705673303586739
InternetService_Fiber optic	0.39369312772910836
Contract_Month-to-month	0.2758324173686368
PhoneService	0.24596280269354162
PaperlessBilling	0.2415529248475717
PaymentMethod_Credit card (automatic)	0.21965729919730578
TechSupport	0.20568534139028977
PaymentMethod_Mailed check	0.2003460678487564
InternetService_DSL	0.19989215054796086
OnlineSecurity	0.19872298622421242
Contract_One year	0.19750966997494915
StreamingMovies	0.17672114393745927
MultipleLines	0.1758751654808676
PaymentMethod_Bank transfer (automatic)	0.1588764197893863
StreamingTV	0.15754563756482862
SeniorCitizen	0.14818334945641423
Dependents	0.10935270769604277
OnlineBackup	0.10888267047331805
gender	0.07031909228841961
DeviceProtection	0.05763341863163266



PaymentMethod_Electronic check	0.057132621248535076
Partner	0.05291432629360901